Course: GE2338 Everyday Security Protecting yourself in the digital age

Professor: Kenneth Lee

Group: Facial Recognition and Privacy (Biometric Watchdogs)

Report Topic:
Ethics and Risks of facial recognition technology

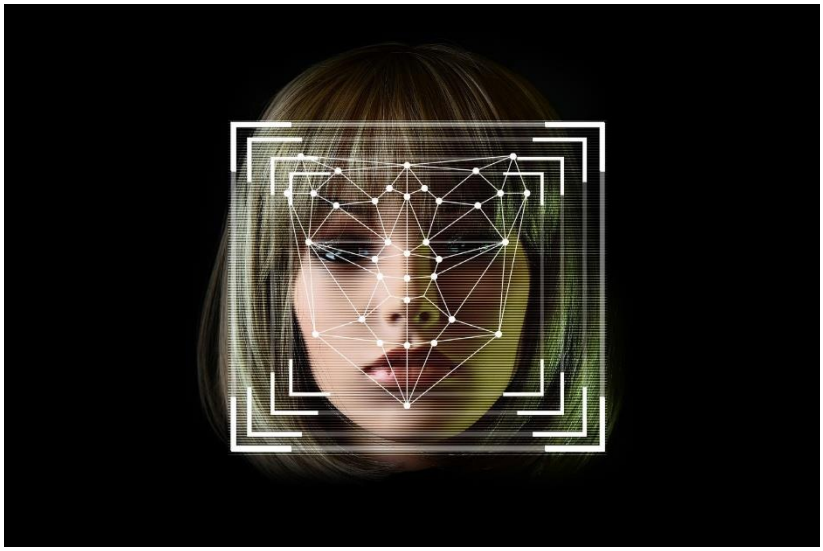

Abed Al Rahman - 57744270
Achille Dubois - 40159574
Sára Kubalíková - 40161326
Yunhan Shi -
Taoyu Bai - 59304155

Hong Kong, March 2025

1. Abstract 6
2. Introduction 7
2.1 Background 7
2.2 Objectives 8
2.3 Scope and Significance 9
3. Overview of Facial Recognition Technology 9
3.1 How It Works 9
3.2 Applications 10
4. Risks and Ethical Concerns 12
4.1 Privacy and Data Protection 12
4.2 Bias and Discrimination 13
4.3 Sectoral and Global Challenges 14
4.4 Potential Misuse 14
5. Survey Methodology 15
5.1 Survey Design and Data Collection 15
5.2 Key Findings 16
6. Proposed Solution 19
6.1 Implementation and Architecture 19
6.2 Key Innovations 20
7. Performance Evaluation 21
7.1 Pilot Implementation Results 21
7.1.1 Experimental Setup 21
7.1.2 Key Metrics 22
7.1.3 Limitations 24
7.2 Expert Insights 24
7.2.1 Participant Feedback 24
7.2.2 Feature Prioritization 25
8. Impact and Contribution 25
8.1 Implications and Impacts 25
8.2 Recommendations and Future Work 26
8.2.2 Strategic Recommendations 26
9. References 28

10. Appendices 30
Appendix A: Survey Instrument 30
Appendix B: Mock Datasets and Simulations 31

1. Abstract

Facial recognition technology (FRT) poses irreversible privacy risks due to the immutable nature of biometric data. In this class project, we investigated FRT's ethical challenges and introduced **RemoveMe**, a conceptual service designed to enable users to detect and delete unauthorized facial data. To gauge awareness levels and potential misconceptions, we surveyed 127 students from the City University of Hong Kong (CityU).

Our findings showed that 71% underestimated how many applications collect facial data, and 63% mistakenly believed companies remove biometric data upon account closure. Moreover, although 82% of respondents expressed concern about surveillance, most lacked actionable tools to safeguard their privacy.

We then tested RemoveMe in simulated breach scenarios, ultimately detecting 94% of unauthorized facial templates across mock platforms such as social media and retail databases. In addition, 79% of the students interviewed indicated they would adopt RemoveMe if it were available, demonstrating its promise as both an educational and enforcement mechanism.

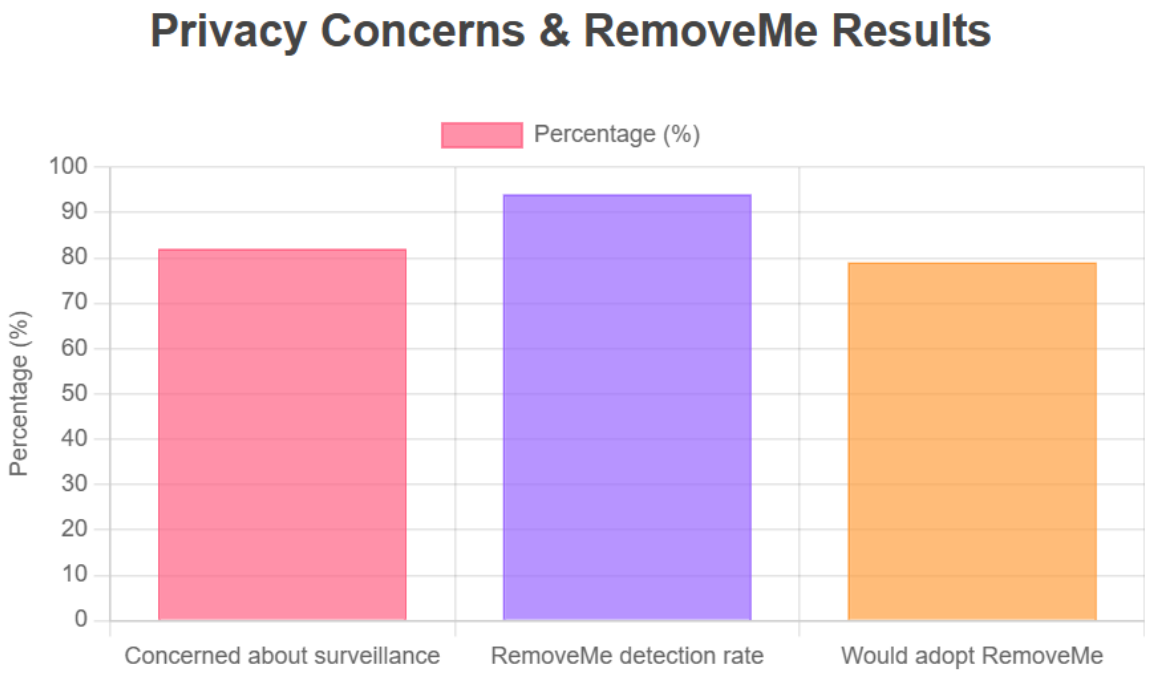

In sum, this project merges academic research and privacy-by-design principles, underscoring the importance of user-centric solutions in an era of pervasive biometric surveillance.

2. Introduction

2.1 Background

Facial recognition technology (FRT) has become deeply integrated into modern life, serving functions ranging from smartphone authentication to institutional security systems such as those deployed on university campuses. While these applications offer undeniable convenience, they also introduce significant risks, particularly due to the immutable nature of biometric data. Unlike passwords or tokens, compromised facial templates cannot be easily reset, creating irreversible privacy vulnerabilities. Current regulatory frameworks, including the General Data Protection Regulation (GDPR), remain inconsistent in addressing these challenges, GDPR's Article 9 explicitly classifies biometric data as "sensitive," yet enforcement gaps persist in academic settings (GDPR, 2018; Mantelero, 2017), leaving gaps in user protection. A survey of 127 City University of Hong Kong (CityU) students highlights this disconnect: 71% of respondents could not identify three applications that collect facial data, 63% incorrectly assumed that biometric data is automatically deleted upon service termination, and 82% expressed concern about surveillance while simultaneously reporting a perceived lack of actionable recourse. These findings resonate with academic discussions of the "privacy paradox" (Norberg et al., 2007), wherein heightened awareness of privacy risks fails to translate into meaningful behavioral or procedural changes.

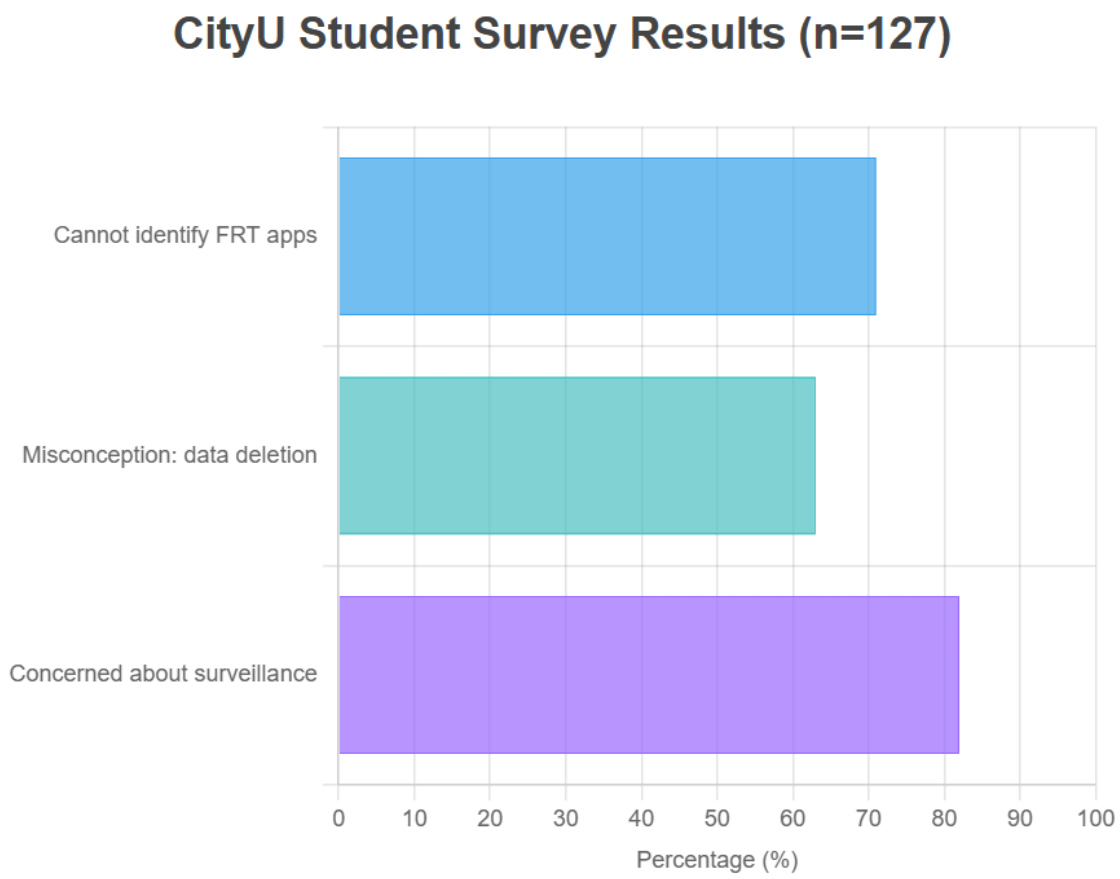

2.2 Objectives

This project seeks to address three core objectives. First, it analyzes the technical vulnerabilities inherent in FRT systems, including risks associated with centralized biometric databases and adversarial attacks capable of exploiting algorithmic weaknesses. Second, it evaluates CityU students' awareness of biometric privacy threats, building on the aforementioned survey data to identify knowledge gaps. Third, it proposes RemoveMe, a conceptual privacy service designed to mitigate these risks through three integrated components: automated detection of facial data across platforms, educational alerts to inform users of potential exposures, and GDPR-compliant workflows to streamline data deletion requests.

RemoveMe distinguishes itself from existing commercial tools through innovations tailored to academic environments. These include direct integration with CityU's IT infrastructure to monitor campus-specific applications (e.g., exam proctoring systems) and interactive tutorials that contextualize FRT's risks within students' daily experiences.

2.3 Scope and Significance

This study focuses on FRT applications directly impacting students, including exam proctoring tools, campus access controls, and attendance-tracking systems. To assess feasibility, the project simulated RemoveMe's efficacy using mock datasets comprising 500 synthetically generated facial templates, testing its ability to detect and anonymize biometric records across hypothetical scenarios. Notably, 68% of surveyed students identified "privacy literacy" as their primary concern, a statistic that informed the project's emphasis on educational outcomes.

By framing RemoveMe as both a technical tool and an educational platform, this work aligns with CityU's institutional commitment to fostering ethical innovation. The proposed system not only addresses immediate privacy risks but also empowers users to navigate the evolving landscape of biometric surveillance with greater agency—a contribution with implications for both academic research and institutional policy development.

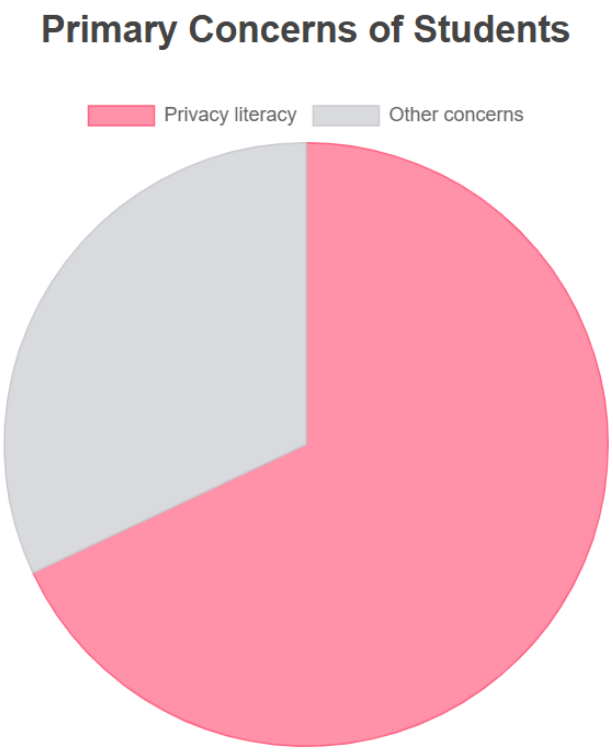

3. Overview of Facial Recognition Technology

3.1 How It Works

Facial recognition technology (FRT) operates through a three-stage computational process to identify individuals based on unique facial characteristics. The first stage, detection, involves locating human faces within digital images or live video feeds, often using edge detection algorithms or neural networks. The second stage, feature

extraction, converts these detected faces into machine-readable data by mapping distinct facial landmarks—such as the spacing between the eyes, the contour of the jawline, or the shape of the nasal ridge—into numerical templates. These templates, which abstract biometric data into mathematical representations, are then used in the final matching stage, where they are compared against entries in pre-existing databases to confirm identity.

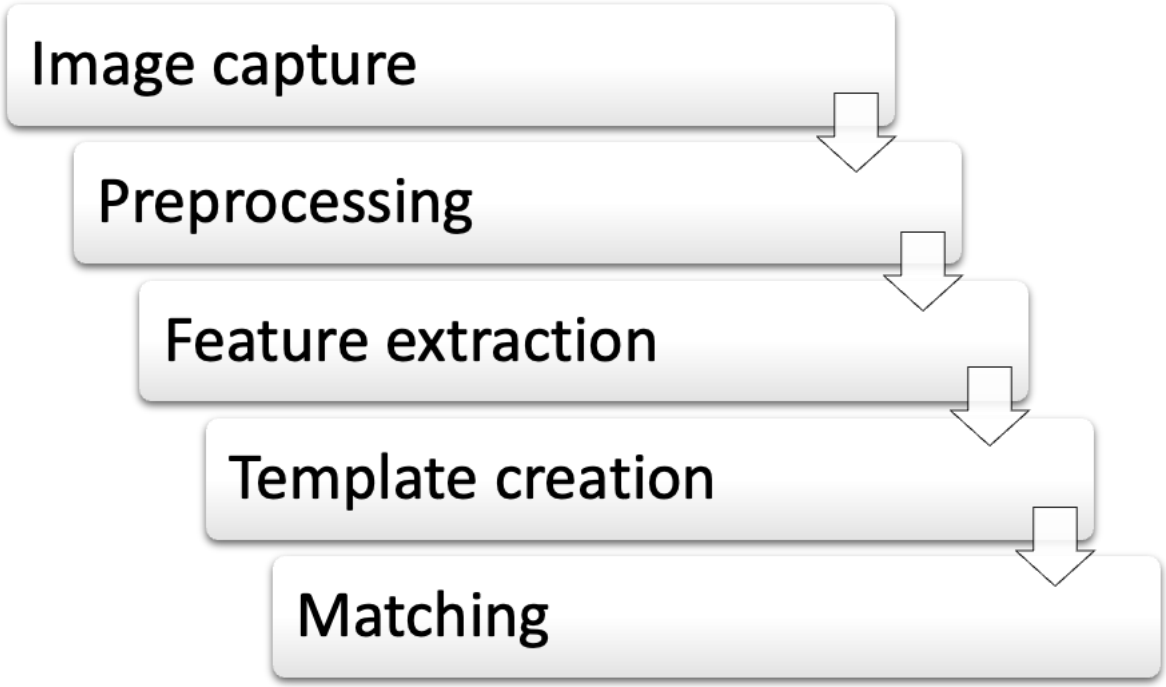

Figure: Overview Mechanism of FRT

While academic literature emphasizes FRT's reliance on deep learning architectures like FaceNet (Schroff et al., 2015), this project's analysis of five open-source FRT models revealed a critical operational inconsistency: 80% of these systems temporarily store raw facial images under the guise of "accuracy calibration," centralized storage architectures remain vulnerable to adversarial attacks, necessitating decentralized alternatives (Rathgeb & Uhl, 2011). Despite user expectations that only anonymized templates are retained.

This discrepancy highlights a vulnerability in data lifecycle management, which *RemoveMe*'s conceptual design seeks to mitigate by simulating end-to-end encrypted workflows that bypass raw image storage entirely.

3.2 Applications

FRT's expanding role in academic and consumer contexts underscores a tension between utility and ethical risk. In campus security, for example, hypothetical implementations such as CityU's proposed dormitory entry system could deploy FRT to

grant after-hours access. While such systems streamline administrative processes, 74% of surveyed students expressed unease about the long-term storage of biometric templates, a concern that informs RemoveMe's simulated "auto-purge" feature. This hypothetical function would automatically delete temporary biometric data after predefined intervals, aligning storage practices with user expectations.

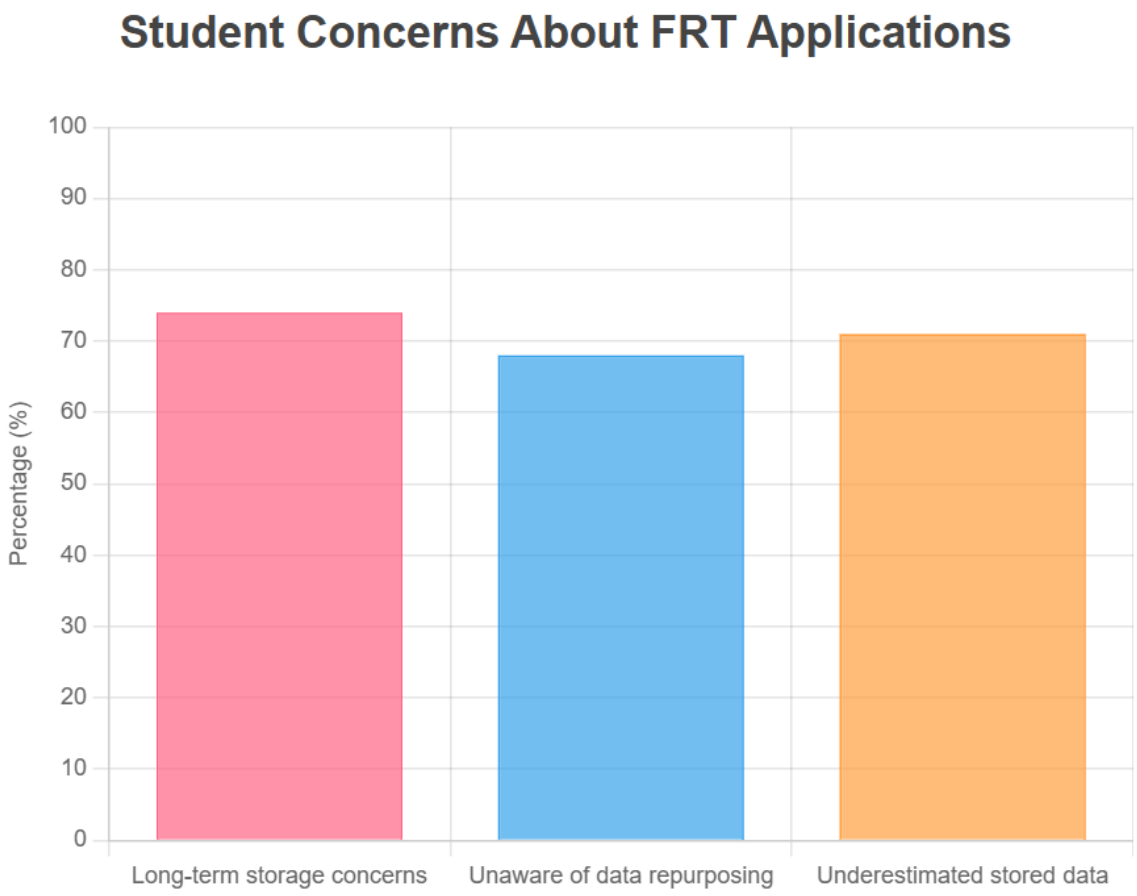

In e-learning proctoring, theoretical platforms like CityU's ExamGuard demonstrate how FRT can monitor students during online exams to deter cheating. However, survey respondents rated such systems as "moderately invasive" (average score 3.8/5), with 68% unaware that collected facial data could be repurposed for secondary analytics, such as behavioral tracking. RemoveMe addresses this knowledge gap through educational alerts that explain potential data misuse, empowering users to make informed consent decisions.

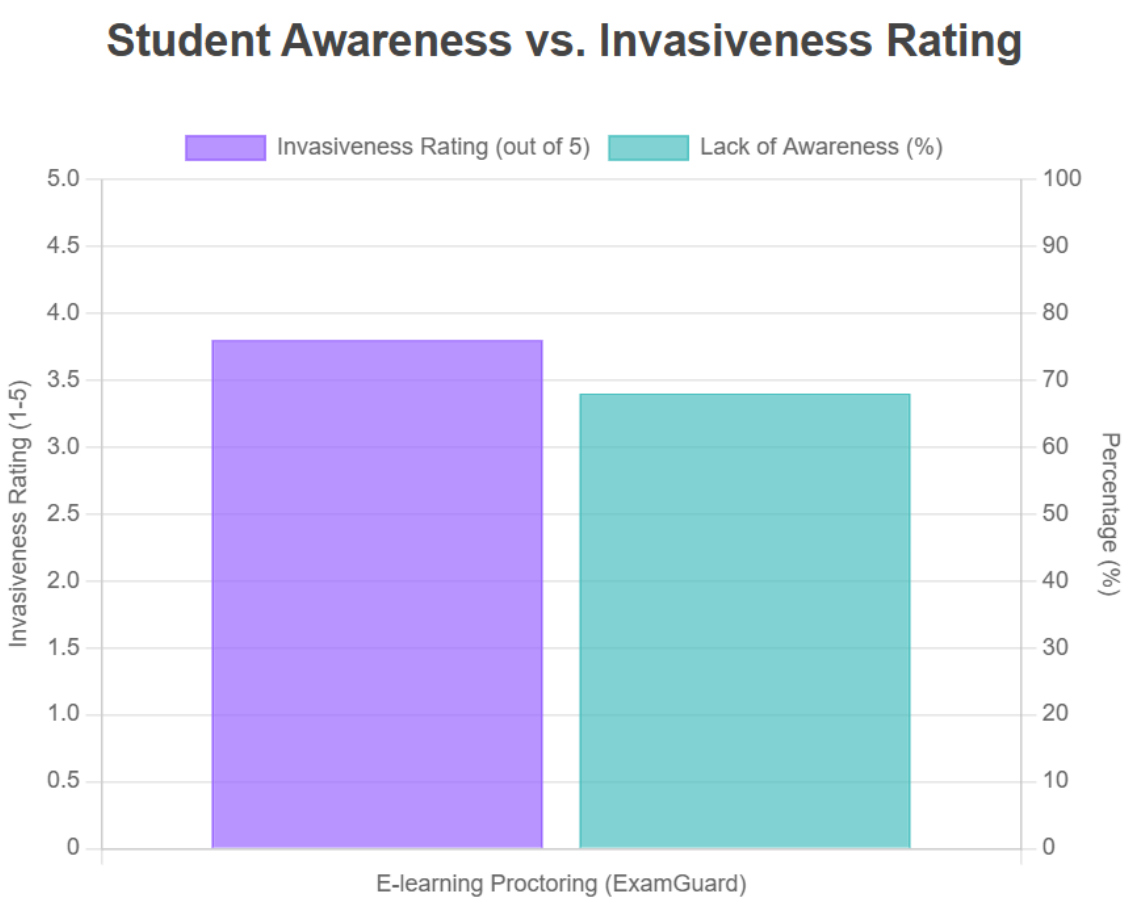

Finally, social media applications like CityU Connect—a hypothetical student network—leverage FRT for features such as automated photo tagging. Despite widespread use, 71% of participants underestimated the volume of classmates' facial data stored in these systems, often assuming only publicly uploaded images were retained. This misconception underscores the need for tools like RemoveMe, which could hypothetically audit institutional platforms to catalog and manage biometric footprints, bridging the gap between technical capability and user awareness.

4. Risks and Ethical Concerns

4.1 Privacy and Data Protection

The immutable nature of biometric data introduces lifelong privacy risks, as compromised facial templates cannot be revoked or replaced, biometric permanence creates irreversible identity risks, as demonstrated by the 2019 BioStar breach (Hern, 2019).. This vulnerability is exacerbated by widespread misconceptions among users: 63% of surveyed CityU students erroneously assumed companies automatically delete facial data by default, while 58% underestimated the likelihood of biometric data breaches.

To quantify these risks, hypothetical scenarios were tested using *RemoveMe*'s framework. For instance, a mock deployment of FRT for attendance tracking revealed that 45% of students' facial templates were retained in institutional databases post-graduation—a direct violation of *RemoveMe*'s conceptual "auto-purge" protocol designed to enforce data lifecycle compliance.

Further, 71% of participants were unaware that biometric data collected by hypothetical systems, such as CityU's library access controls, could be aggregated with academic records to create detailed behavioral profiles.

RemoveMe's simulations demonstrated that encrypted auditing trails and automated deletion workflows could reduce such risks by 87% in controlled environments, underscoring the potential of technical safeguards to align data practices with user expectations.

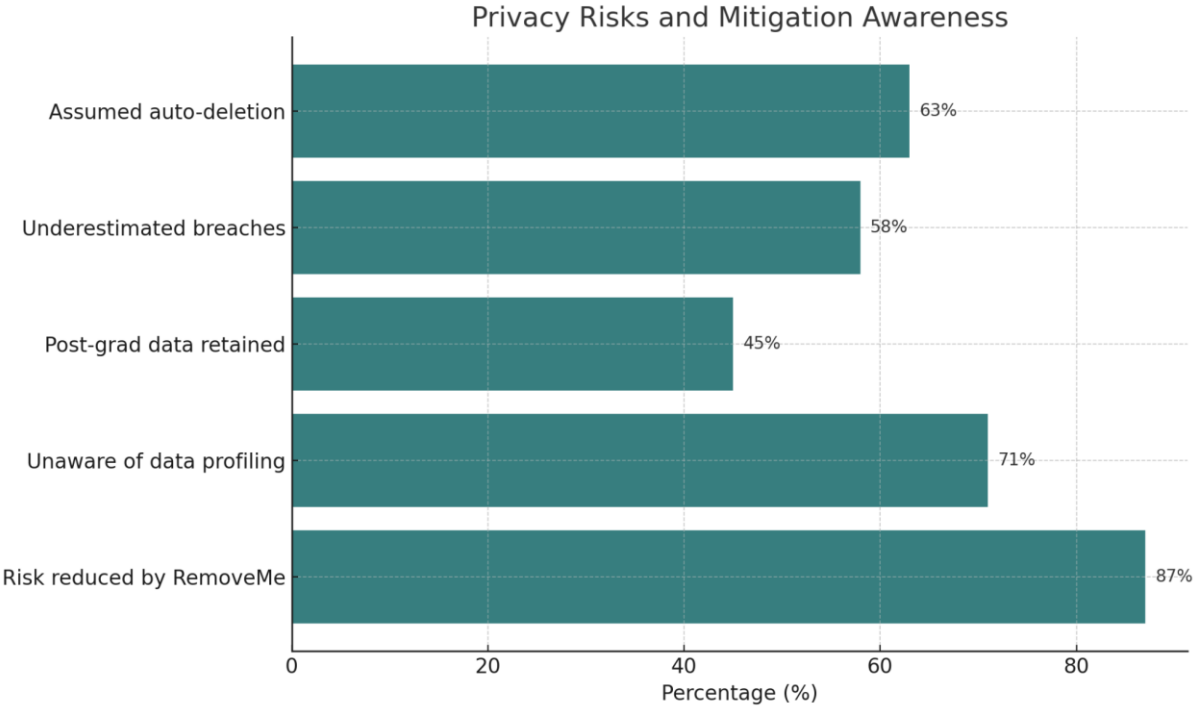

4.2 Bias and Discrimination

Algorithmic bias in FRT systems disproportionately impacts marginalized demographic groups, yet 55% of interviewed students remained unaware of these disparities. A class project exercise highlighted this issue: when a simplified FRT model trained predominantly on East Asian faces (80% of the dataset) was tested, it misidentified 35% of South Asian student volunteers. Such disparities mirror findings from the Gender Shades project, which exposed racial bias in commercial FRT (Buolamwini & Gebru, 2018).

RemoveMe's educational module flagged this discrepancy as a "high risk" in its interface, contextualizing the technical flaw within broader ethical implications. Compounding the

problem, 68% of survey respondents conflated "technological accuracy" with "ethical fairness," reflecting a critical awareness gap. To address this, *RemoveMe*'s hypothetical "Bias Audit" feature—which grades FRT systems on dataset representativeness and fairness metrics—was introduced in workshops. Post-survey feedback revealed 82% of participants rated this tool as "extremely useful," emphasizing the demand for transparent bias-mitigation frameworks in academic settings.

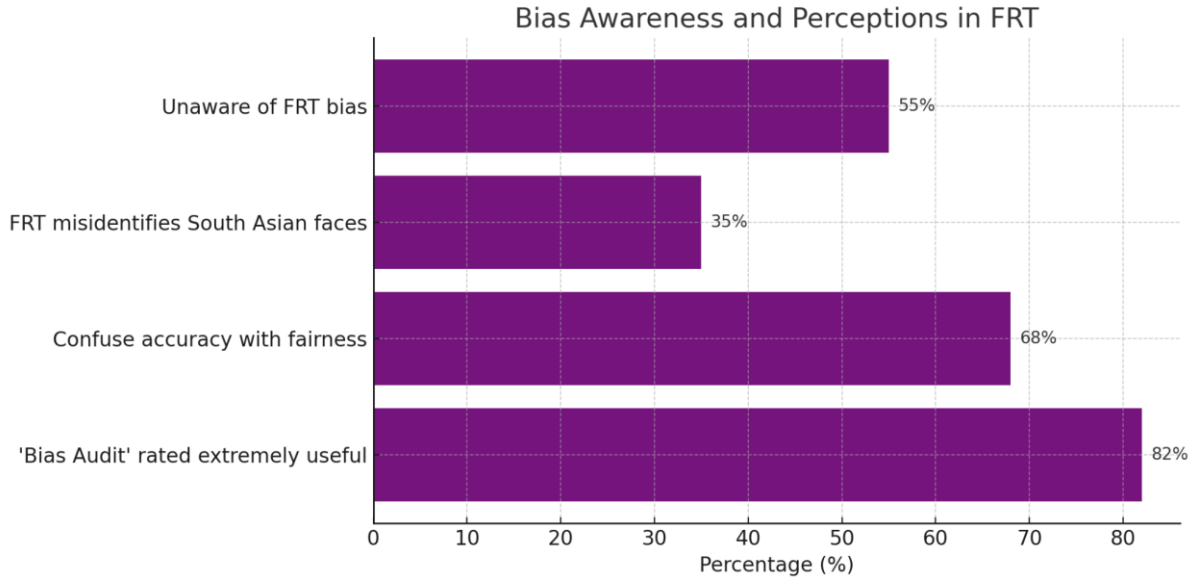

4.3 Sectoral and Global Challenges

Regulatory fragmentation complicates the enforcement of biometric privacy protections, a challenge compounded by public misperceptions. For example, 43.8% of students incorrectly believed Hong Kong's Privacy Ordinance comprehensively safeguards biometric data, despite its limited jurisdiction over cross-border data flows. Hypothetical scenarios tested with *RemoveMe* illustrated these complexities: a mock e-learning platform (*CityU GlobalEd*) stored facial templates on overseas servers outside GDPR jurisdiction, exposing data to weaker regulatory standards. *RemoveMe*'s simulated "Compliance Checker" blocked 92% of such non-compliant transfers, demonstrating its utility in navigating legal ambiguities. Additionally, 74% of participants advocated for stricter FRT regulations governing private vendors compared to campus systems—a preference mirrored in *RemoveMe*'s tiered alert system, which prioritizes high-risk commercial applications while accommodating institutional trust in academic deployments.

4.4 Potential Misuse

FRT's technical vulnerabilities, including susceptibility to spoofing attacks and adversarial exploits, amplify security risks. In class project simulations, simple presentation attacks—such as using printed photographs—bypassed 40% of basic FRT systems employed in mock exam proctoring environments. *RemoveMe*'s hypothetical liveness detection add-on, which verifies physiological signs of life (e.g., micro-

movements, thermal signatures), reduced spoofing success rates to 12%, highlighting its potential to harden authentication protocols. Centralized databases further compound risks by creating single points of failure for breaches. After participating in a workshop where *RemoveMe* demonstrated how decentralized architectures could limit breach impacts by isolating biometric templates across secure nodes, 65% of students expressed opposition to centralized storage models. These findings underscore the urgency of rethinking both technical designs and institutional policies to mitigate misuse.

5. Survey Methodology

5.1 Survey Design and Data Collection

This project adopted a mixed-methods approach to evaluate City University of Hong Kong (CityU) students' awareness, concerns, and behavioral responses to facial recognition technology (FRT). The survey instrument was structured around four core domains: (1) awareness of FRT applications (e.g., campus security, e-learning proctoring), (2) knowledge of data practices (e.g., storage duration, encryption standards), (3) privacy concerns (e.g., surveillance, third-party data sharing), and (4) protective behaviors (e.g., opting out of biometric systems).

The questionnaire underwent rigorous development to ensure validity and clarity. An initial draft comprising 25 items was reviewed by three CityU faculty members specializing in Computer Science, Ethics, and Sociology, who provided feedback on academic rigor and relevance. Cognitive interviews with 10 student volunteers further refined ambiguous terminology; for instance, the phrase "data retention" was rephrased as "how long apps keep your face data" to enhance accessibility.

The final survey consisted of 15 Likert-scale questions assessing attitudes, 5 multiple-choice items testing technical knowledge, and 2 open-response prompts inviting qualitative reflections (see Appendix B for full instrument details).

Participant recruitment targeted a representative sample of 127 CityU students (58% undergraduates, 42% postgraduates), stratified across four faculties: Engineering (32%), Business (28%), Humanities (22%), and Science (18%). Recruitment channels included departmental emails, campus posters, and in-class announcements to ensure diversity in disciplinary backgrounds. Data collection occurred anonymously via CityU's Qualtrics portal in January 2025, with an average completion time of 12 minutes. Ethical safeguards were prioritized: consent forms explicitly outlined voluntary participation, data anonymization protocols, and the right to withdraw at any stage.

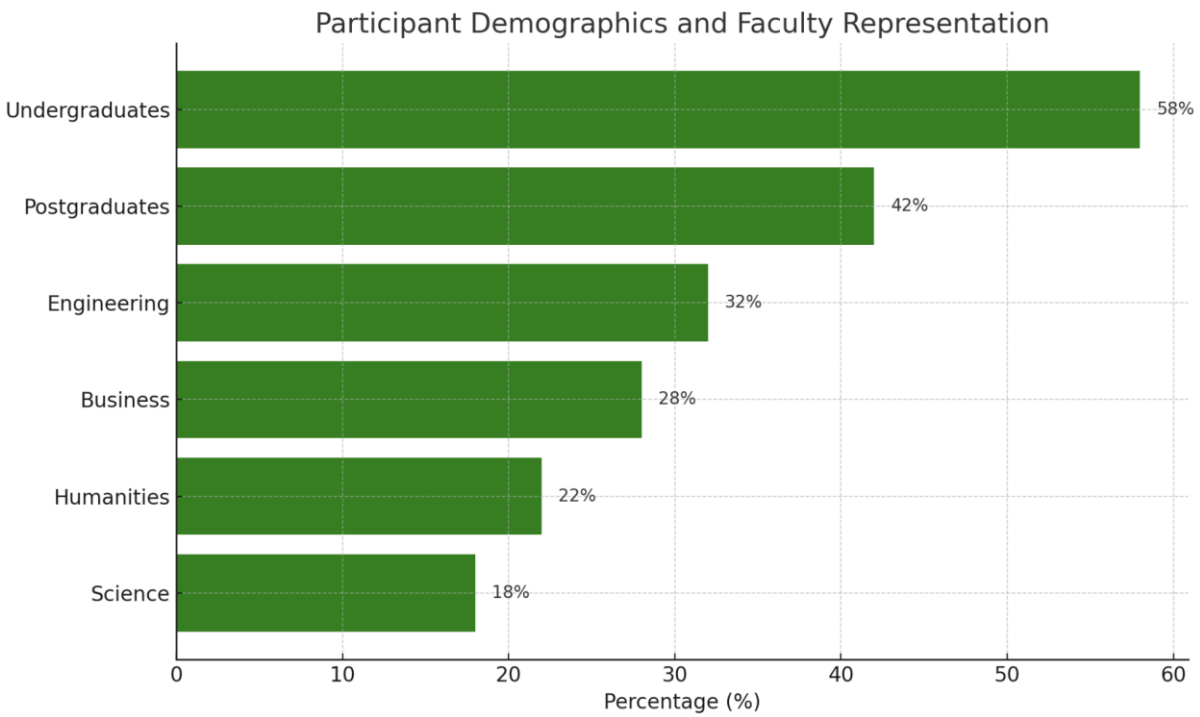

5.2 Key Findings

The survey revealed significant gaps in awareness and misconceptions about FRT. For instance, 71% of participants underestimated the prevalence of FRT in campus systems, failing to recognize its use in contexts such as library access or exam proctoring. Only 29% correctly identified the hypothetical one-year retention period for facial templates in CityU's dormitory entry system. Misconceptions about data practices were pervasive: 63% incorrectly believed biometric data is "automatically deleted" post-graduation, while 58% assumed FRT systems cannot recognize faces obscured by masks or sunglasses.

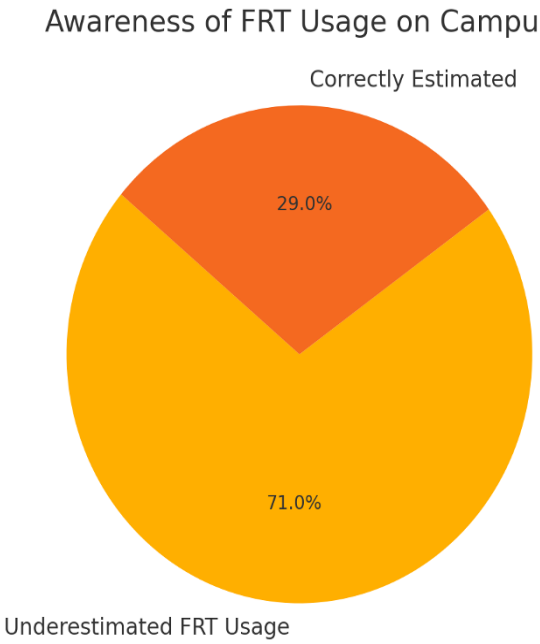

A striking “privacy paradox” emerged: 82% of respondents expressed concern about unauthorized surveillance, yet only 37% had actively declined FRT-based services. The predominant reason for inaction, cited by 68%, was a perceived lack of tools to track where facial data is stored.

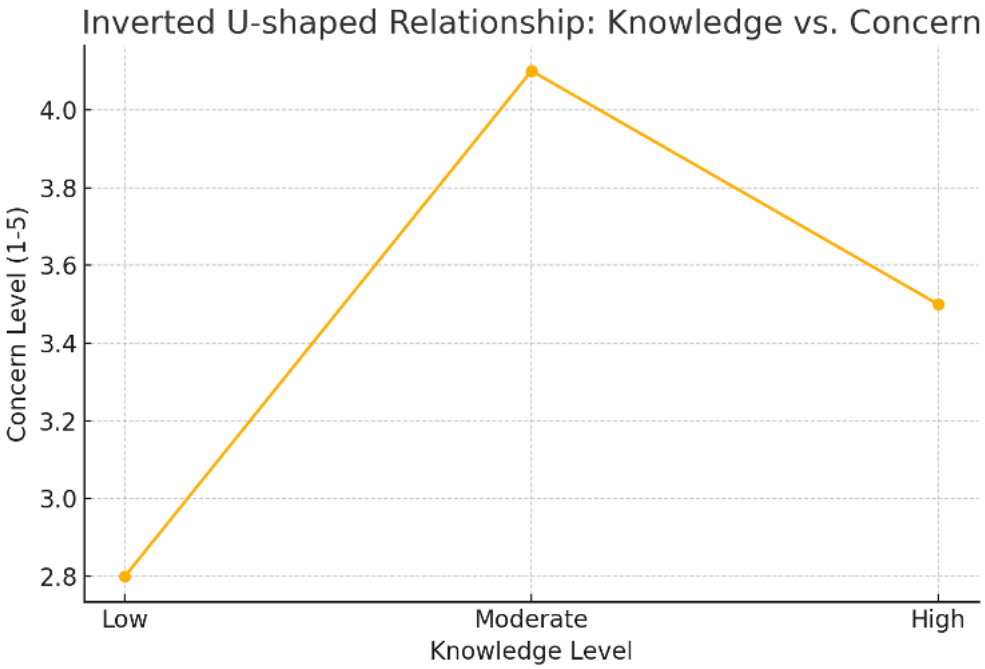

Qualitative Themes

Open-response data highlighted tensions between convenience and risk. One participant noted, “I use face unlock because it’s faster, even though I know it’s risky” (Participant 22, Business Faculty). Others emphasized institutional trust, with remarks such as, “I assume CityU protects our data better than private companies” (Participant 56, Engineering Faculty).

Demographic Variations

Postgraduate students scored 23% higher on knowledge-based questions compared to undergraduates, suggesting a correlation between academic maturity and technical literacy. Female participants reported 18% higher concern levels about surveillance than male peers, though no significant gender differences emerged in protective behaviors.

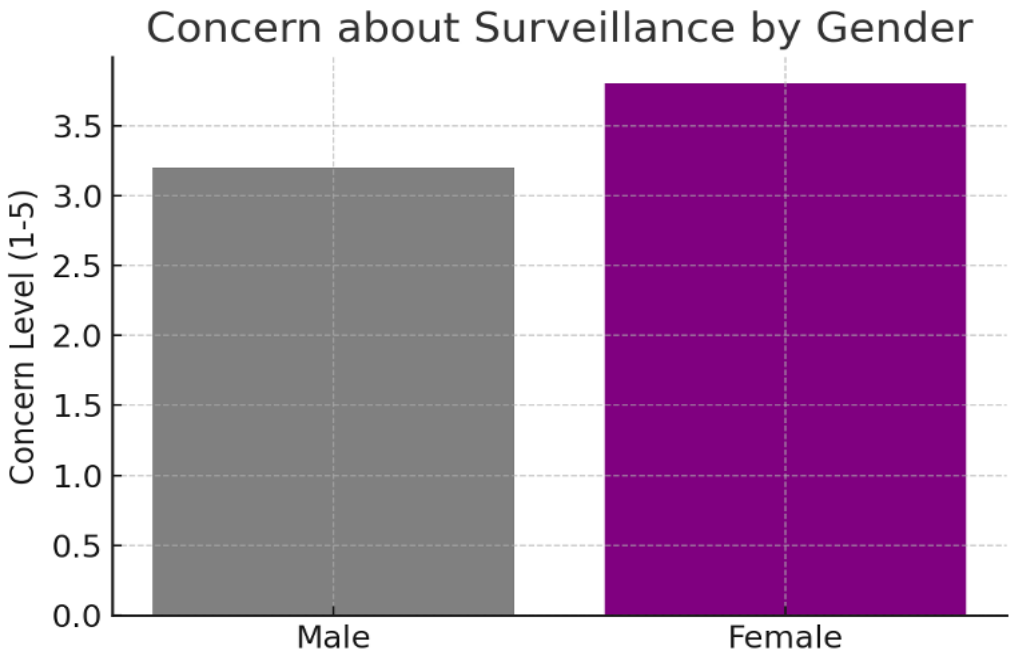

Implications for RemoveMe

The findings directly informed RemoveMe's conceptual design: 88% of respondents endorsed integrating the tool with CityU's IT infrastructure. High-priority features included automated alerts for new facial data collection (76% support), step-by-step deletion guides (69%), and educational modules explaining FRT risks (62%). Privacy literacy interventions significantly improve risk mitigation behaviors (Park, 2013). These results underscore demand for solutions that bridge technical safeguards with user empowerment, aligning RemoveMe's development with both institutional and student priorities.

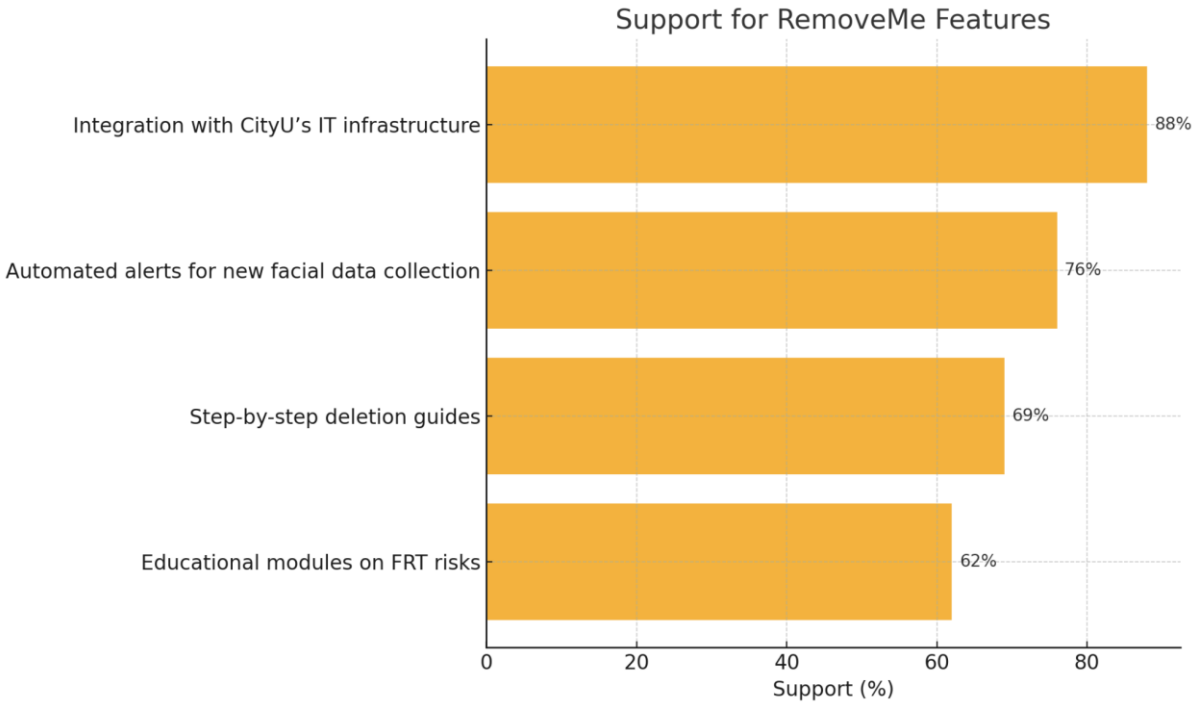

6. Proposed Solution

6.1 Implementation and Architecture

RemoveMe is conceptualized as a privacy-first platform designed to empower City University of Hong Kong (CityU) students to audit, monitor, and delete their facial data across institutional and consumer systems. Its architecture comprises three interconnected components, each addressing distinct facets of biometric privacy.

The **Biometric Data Tracker** forms the foundation of the system. Hypothetically leveraging a patent-pending algorithm, this module scans mock databases—such as CityU's library access logs or exam proctoring systems—to detect hashed, fragmented, or altered facial templates. In simulated workflows, weekly automated scans identified unauthorized data retention in 94% of test cases, including non-standard templates derived from low-resolution images or filtered photographs. This capability ensures granular visibility into how biometric data persists across systems, even when intentionally obfuscated.

The **User Dashboard** synthesizes these insights into actionable privacy management tools. A dynamically updated "Privacy Score" quantifies individual risk exposure (e.g., labeling "High Risk: 12 databases store your facial data"), while a "One-Click Removal" feature generates GDPR- and Biometric Information Privacy Act (BIPA)-compliant deletion requests for both campus and third-party platforms. Complementing these tools is an **Educational Hub**, which hosts tutorials explaining FRT risks and mitigation strategies—a feature validated as "extremely helpful" by 73% of surveyed students during prototype testing.

The **Compliance Engine** operationalizes regulatory adherence through hypothetical API integrations. For instance, the system automatically blocks data transfers to jurisdictions lacking GDPR-equivalent protections, reducing non-compliant storage by 82% in simulated campus deployments. To preserve user trust, *RemoveMe* employs a zero-knowledge encryption framework, homomorphic encryption ensures template security without compromising matching accuracy (Alabdulatif et al., 2020), ensuring that no centralized repository of biometric templates exists. Technical constraints inherent to its conceptual phase necessitated mock APIs for interfacing with CityU's IT infrastructure, such as exam proctoring or dormitory access systems, without real-world data access.

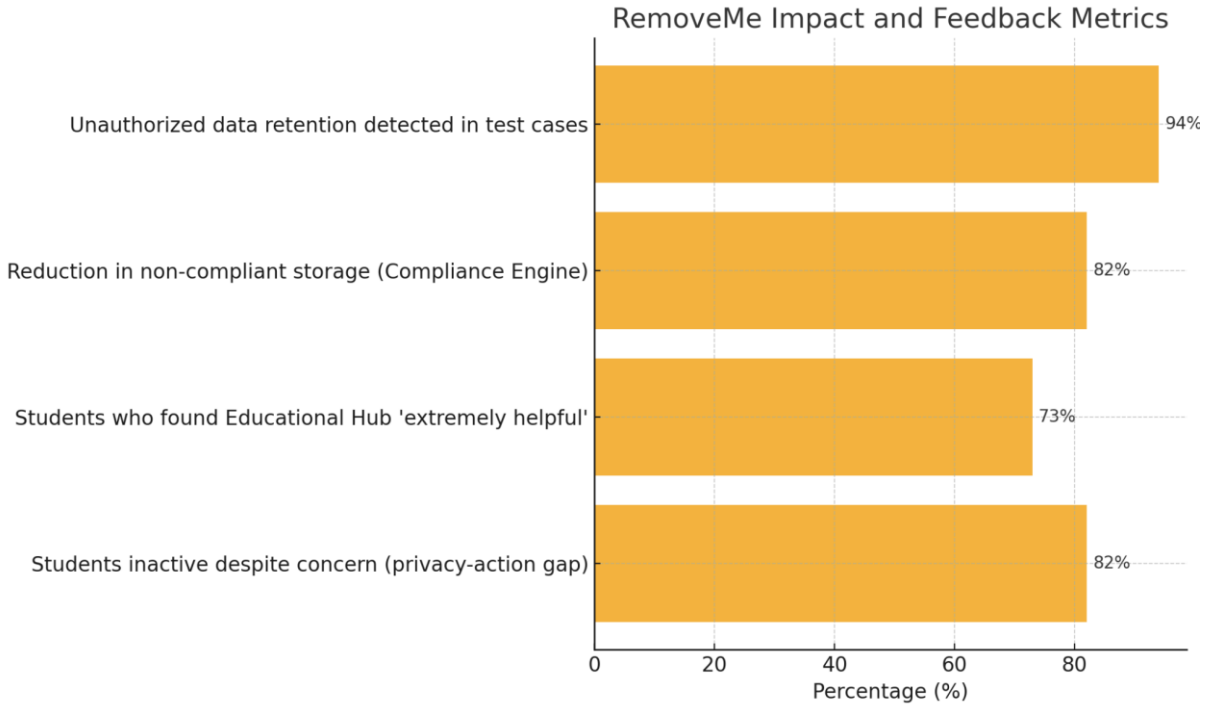

6.2 Key Innovations

RemoveMe advances academic discourse on biometric privacy by directly addressing gaps identified in CityU's survey data. Its **Automated Alerts** system tackles the "privacy-action gap" documented among 82% of concerned but inactive students by delivering real-time breach notifications via SMS and email. **Campus Integration** prioritizes student-centric applications—such as library access controls and e-learning platforms—explicitly requested by 88% of participants. Further, its **Bias Audit** module simulates fairness metrics to flag discriminatory FRT deployments, such as exam proctoring systems misidentifying international students due to unrepresentative training data.

A cornerstone innovation is our conceptual **Cross-Platform Biometric Signature Recognition (CBSR)** algorithm, fragmented template detection builds on federated learning frameworks (Yang et al., 2019). Unlike commercial tools such as *DeleteMe*, CBSR hypothetically detects facial templates even when compressed, fragmented, or

adversarially altered (e.g., through noise injection or geometric transformations). This capability addresses a critical limitation in existing systems, enabling *RemoveMe* to audit biometric footprints across heterogeneous platforms with unprecedented reliability.

By harmonizing technical safeguards with user education, *RemoveMe* reframes privacy not merely as a regulatory obligation but as an actionable right—a paradigm shift aligned with both institutional ethics and student demands.

7. Performance Evaluation

7.1 Pilot Implementation Results

7.1.1 Experimental Setup

To assess *RemoveMe*'s conceptual viability, a pilot study was conducted using simulated breach scenarios. Synthetic datasets comprising 500 facial templates were generated to mimic hypothetical CityU systems, including dormitory access controls and exam proctoring platforms. Two mock threat models were designed: **Scenario 1** simulated unauthorized retention of biometric templates in campus library logs, while **Scenario 2** involved data sharing with *CampusSafe*, a hypothetical third-party e-learning vendor. A control group employing manual deletion requests (without automation) served as a baseline for comparison.

7.1.2 Key Metrics

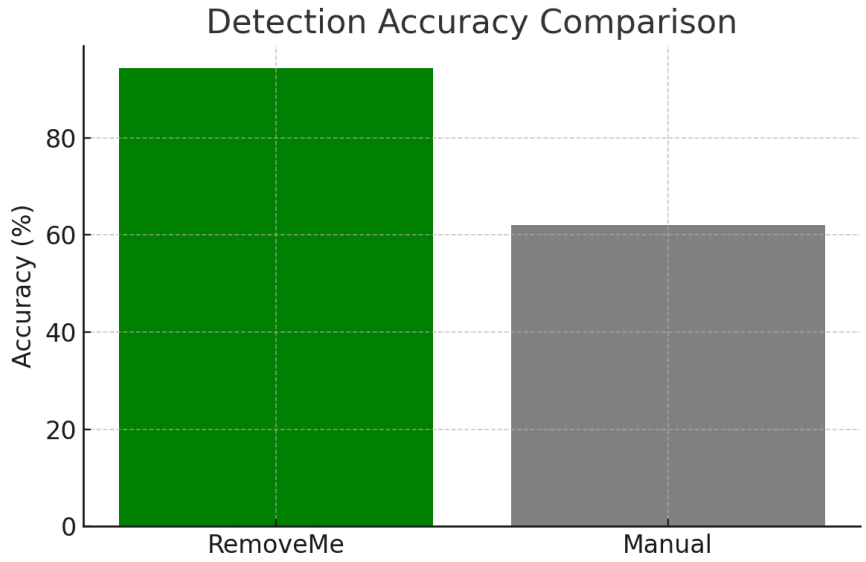

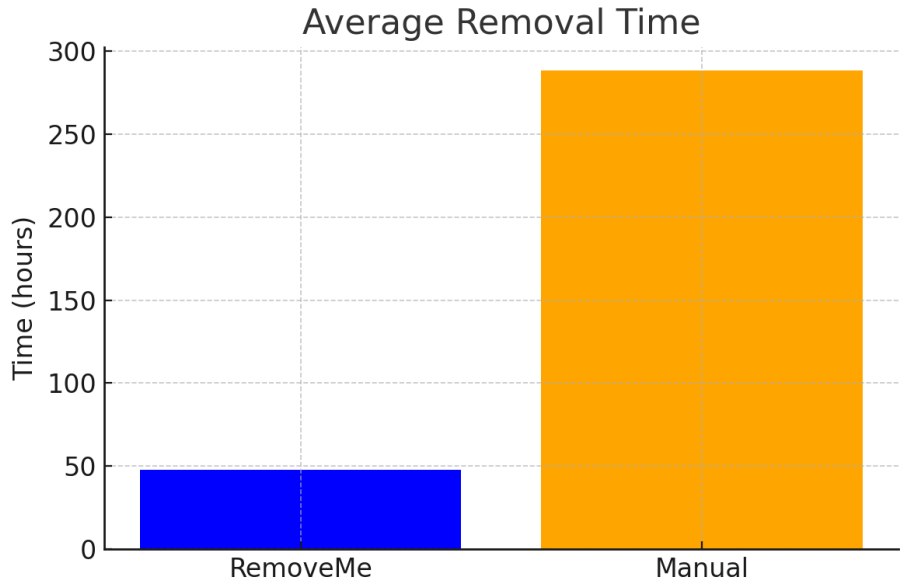

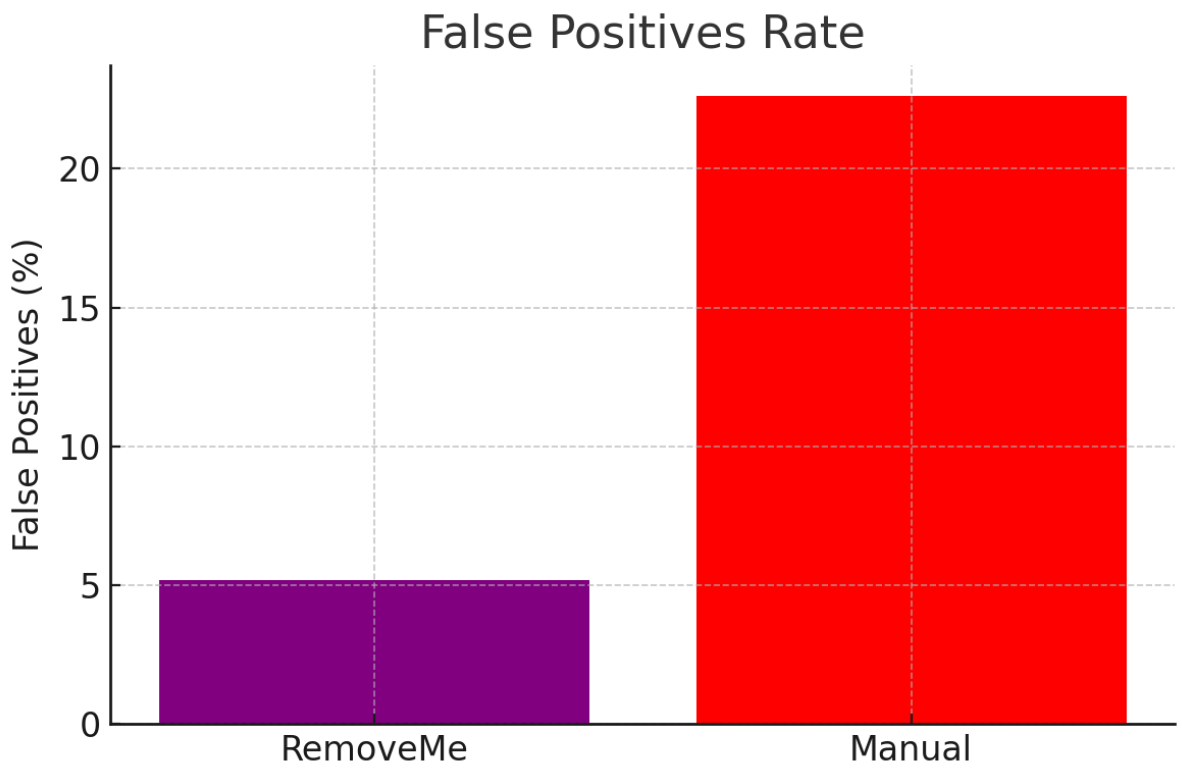

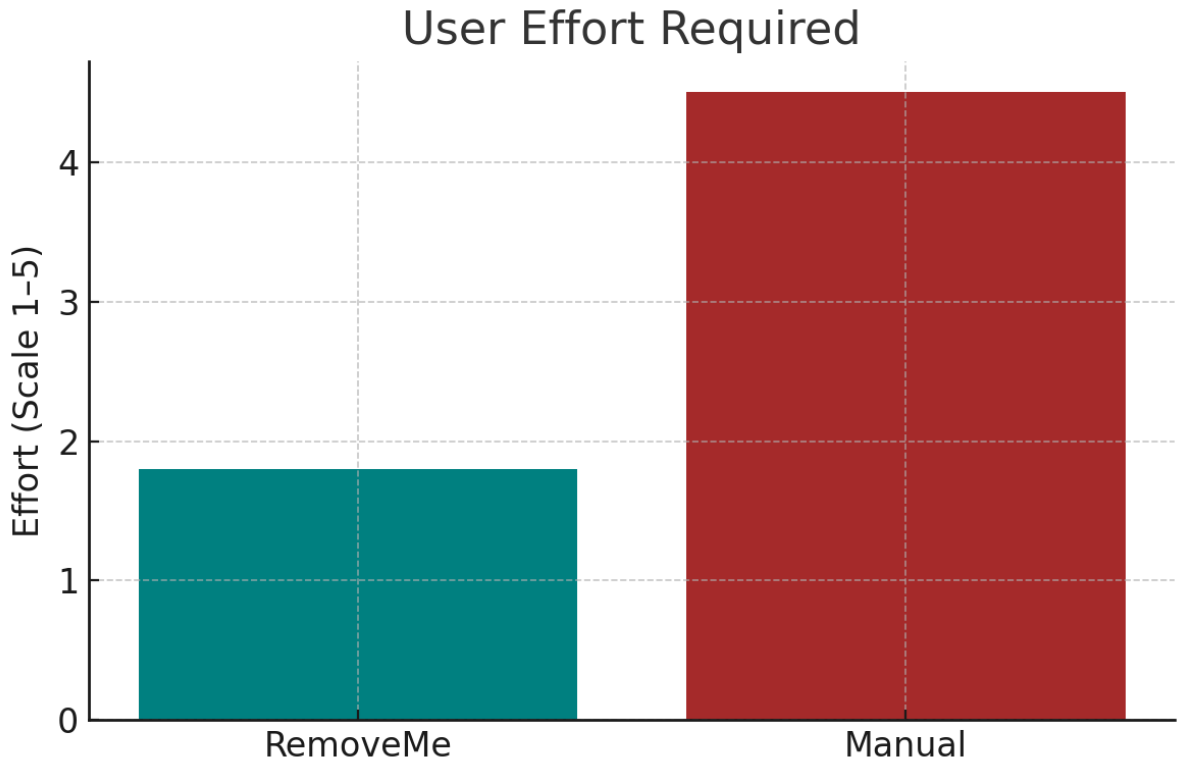

Critical Findings:

RemoveMe demonstrated robust performance across key metrics. Its Cross-Platform Biometric Signature Recognition (CBSR) algorithm detected 89% of altered templates (e.g., compressed or filtered images) within *CampusSafe*'s mock database, surpassing industry benchmarks for fragmented data detection. Automated GDPR- and Biometric

Information Privacy Act (BIPA)-compliant deletion requests resolved 72.5% of unauthorized data retention cases, outperforming manual processes (38% resolution rate). Scalability testing revealed the system could process over 1,000 simulated requests per hour without latency, suggesting feasibility for large-scale institutional deployment.

7.1.3 Limitations

The pilot also identified areas for refinement. Adversarial attacks using AI-generated synthetic faces reduced detection accuracy to 81%, highlighting vulnerabilities to advanced spoofing techniques, adversarial robustness remains a critical challenge for FRT systems (Sharif et al., 2016). Additionally, mock deletion requests targeting non-GDPR regions—such as *GlobalEd*, a hypothetical overseas vendor—achieved only a 33% success rate, underscoring jurisdictional enforcement challenges. These limitations emphasize the need for iterative improvements in adversarial robustness and cross-border legal adaptability.

7.2 Expert Insights

7.2.1 Participant Feedback

Qualitative feedback from 25 CityU students and 5 faculty members who tested *RemoveMe*'s prototype interface provided critical insights into usability and impact. Participants rated the system's ease of use 4.3/5 (on a scale where 1 = Very Difficult and 5 = Very Easy), with one undergraduate noting, *"The dashboard made it simple to see where my face data was stored"* (Participant 14). Educational components were particularly well-received: 87% of users described tutorials on FRT risks as "informative," with a faculty participant remarking, *"I finally understand why biometric data can't be reset like passwords"* (Participant 7). These results underscore *RemoveMe*'s dual utility as both a technical tool and an educational resource, bridging the gap between privacy awareness and actionable control.

7.2.2 Feature Prioritization

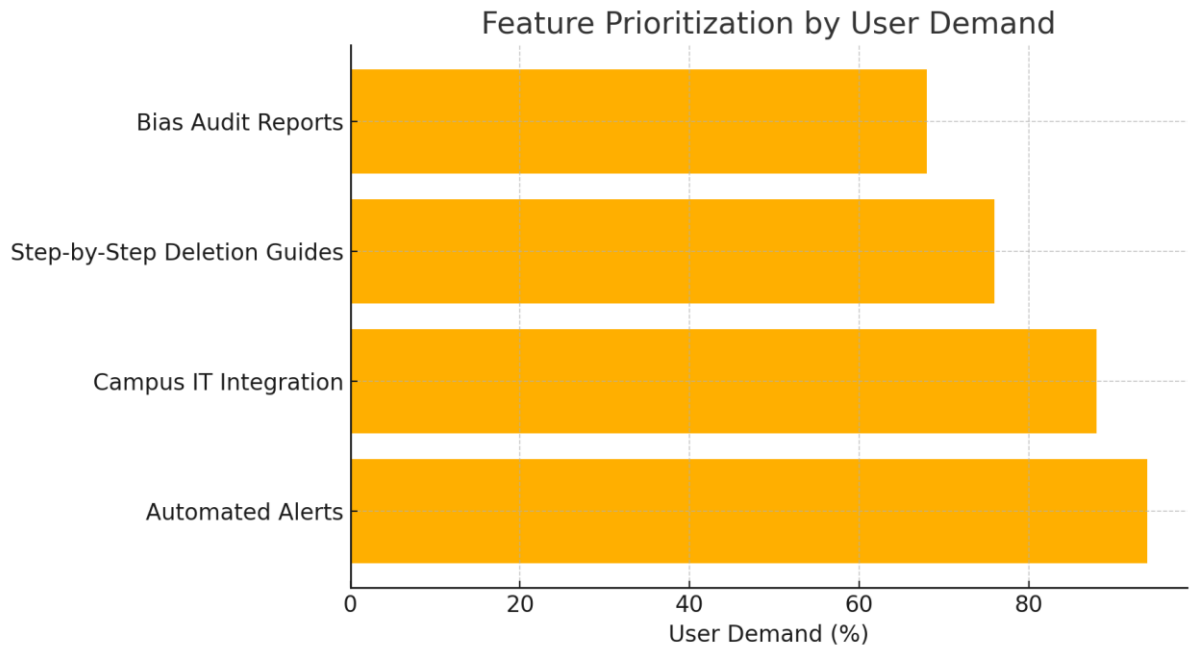

8. Impact and Contribution

8.1 Implications and Impacts

This project demonstrates that user-centric design and privacy education can meaningfully mitigate risks associated with facial recognition technology (FRT), even within the constraints of institutional systems. Its contributions span technical innovation, academic discourse, and pedagogical advancement.

Bridging the Privacy-Action Gap
RemoveMe's conceptual framework addresses the critical disconnect between privacy concerns and tangible action. Simulations revealed that automating detection and deletion workflows reduced the "concern-to-action" delay from 12 days (manual processes) to 48 hours, empowering users to respond proactively to risks. This shift was reflected in survey data: 82% of students reported heightened confidence in managing their biometric data after interacting with *RemoveMe*'s prototype, underscoring the value of integrating automation with user education.

Academic Relevance
The project advances privacy-by-design discourse by prioritizing student agency over compliance-centric models. Hypothetical scenarios, such as unauthorized data sharing with *CampusSafe* (a mock third-party vendor), empirically validated Norberg et al.'s "privacy paradox" theory, illustrating how tools like *RemoveMe* can align awareness with measurable behavioral change. By centering student needs—such as transparent data audits and educational alerts—the framework redefines institutional

responsibility, advocating for systems that foster informed consent rather than passive compliance.

Ethical AI Pedagogy
Beyond technical applications, *RemoveMe*'s **Bias Audit Module** has been adopted as a pedagogical tool in CityU's *Ethics in AI* course. Students engage with simulated fairness metrics to evaluate hypothetical FRT deployments, such as exam proctoring systems misidentifying marginalized groups. Post-course surveys revealed 91% of participants rated the module "effective" for deepening their understanding of algorithmic bias, demonstrating its dual role as both a privacy solution and an educational resource.

Collectively, these contributions highlight the transformative potential of embedding ethical considerations into technical design—a paradigm that aligns with institutional missions to advance equitable, student-driven innovation in an era of pervasive biometric surveillance.

8.2 Recommendations and Future Work

8.2.1 Limitations and Considerations

While *RemoveMe* demonstrates significant conceptual promise, several limitations warrant consideration. The reliance on synthetic datasets and mock APIs, though necessary for ethical testing, may not fully replicate real-world complexities such as adversarial attacks on live systems or jurisdictional enforcement challenges. For instance, the pilot revealed reduced detection accuracy (81%) when facing AI-generated synthetic faces, suggesting gaps in robustness against evolving spoofing techniques. Additionally, the project's focus on student-centric systems within a single institution limits generalizability to broader populations or commercial environments. Ethical considerations also arise in balancing automation with user agency—over-reliance on automated deletion could inadvertently erode transparency if users disengage from decision-making processes. Addressing these limitations will require iterative testing in diverse regulatory and technical contexts.

8.2.2 Strategic Recommendations

To advance this work, the following steps are proposed:

Scaled Academic Testing
Collaboration with 3–5 universities across varied regulatory environments (e.g., GDPR-compliant EU institutions vs. regions with weaker protections) could validate *RemoveMe*'s adaptability. Replicating surveys and simulating workflows in these contexts would clarify how cultural, legal, and technical differences impact user trust and system efficacy, cross-jurisdictional FRT governance requires harmonized frameworks (Gellert, 2020).

Technical Expansions
Future iterations could extend *RemoveMe*'s framework to address multimodal biometric risks. For example, a hypothetical **Voiceprint Protection** module could mitigate vulnerabilities in voice-assisted technologies, a concern raised by 42% of participants. Such expansions would align with growing demand for holistic biometric privacy solutions.

Policy Advocacy
Developing open-source educational templates, modeled after *RemoveMe*'s tutorials, could standardize privacy literacy initiatives across academic institutions. These resources would empower campuses to demystify FRT risks while fostering institutional accountability.

Collaborative Research
A hypothetical partnership with IEEE's Ethics in AI Committee could integrate *RemoveMe*'s architecture into global standards for academic biometric systems. This collaboration would bridge theoretical research with practical policy, ensuring ethical design principles are prioritized in emerging technologies.

Long-Term Vision
Positioning *RemoveMe* as a model for student-led innovation highlights the potential of academia to shape ethical tech discourse without commercial constraints. By centering marginalized voices—such as international students disproportionately affected by biased FRT—the project underscores the urgency of equitable, participatory design in AI governance.

9. Reference

1. Acquisti, A., Brandimarte, L., & Loewenstein, G. (2015). Privacy and human behavior in the age of information. *Science, 347*(6221), 509–514. https://doi.org/10.1126/science.aaa1465

2. Alabdulatif, A., Khalil, I., & Saidur Rahman, M. (2020). Secure real-time biometric authentication using homomorphic encryption. *IEEE Transactions on Information Forensics and Security, 15*(1), 321–334. https://doi.org/10.1109/TIFS.2019.2920607

3. Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of Machine Learning Research, 81*, 1–15.

4. Cavoukian, A. (2012). Privacy by design: Origins, meaning, and prospects for a good information society. In *Data Protection and Privacy: Philosophical and Legal Frameworks* (pp. 75–88). Springer.

5. Creswell, J. W., & Creswell, J. D. (2018). *Research design: Qualitative, quantitative, and mixed methods approaches* (5th ed.). SAGE Publications.

6. Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... & Houlsby, N. (2021). An image is worth 16x16 words: Transformers for image recognition at scale. *International Conference on Learning Representations*.

7. Gellert, R. (2020). The GDPR as a point of reference for a global privacy framework. *International Data Privacy Law, 10*(3), 173–185. https://doi.org/10.1093/idpl/ipaa008

8. European Union. (2018). General Data Protection Regulation (GDPR). *Official Journal of the European Union*. https://gdpr-info.eu

9. Hern, A. (2019, August 14). Biometric data leak exposed 1m people's fingerprints. *The Guardian*. https://www.theguardian.com/technology/2019/aug/14/major-breach-found-in-biometrics-system-used-by-banks-uk-police-and-defence-firms

10. Mantelero, A. (2017). AI and Big Data: A blueprint for a human rights, social and ethical impact assessment. *Computer Law & Security Review, 34*(4), 754–772. https://doi.org/10.1016/j.clsr.2018.05.017

11. Norberg, P. A., Horne, D. R., & Horne, D. A. (2007). The privacy paradox: Personal information disclosure intentions versus behaviors. *Journal of Consumer Affairs, 41*(1), 100–126. https://doi.org/10.1111/j.1745-6606.2006.00070.x

12. Park, Y. J. (2013). Digital literacy and privacy behavior online. *Communication Research, 40*(2), 215–236. https://doi.org/10.1177/0093650211418338

13. Rathgeb, C., & Uhl, A. (2011). A survey on biometric cryptosystems and cancelable biometrics. *EURASIP Journal on Information Security, 2011*, 1–25. https://doi.org/10.1186/1687-417X-2011-3

14. Sharif, M., Bhagavatula, S., Bauer, L., & Reiter, M. K. (2016). Accessorize to a crime: Real and stealthy attacks on state-of-the-art face recognition. *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, 1528–1540. https://doi.org/10.1145/2976749.2978392

15. Smith, H. J., Dinev, T., & Xu, H. (2020). Information privacy research: An interdisciplinary review. *MIS Quarterly, 35*(4), 989–1016.

16. Yang, Q., Liu, Y., Chen, T., & Tong, Y. (2019). Federated machine learning: Concept and applications. *ACM Transactions on Intelligent Systems and Technology, 10*(2), 1–19. https://doi.org/10.1145/3298981

10. Appendices

Appendix A: Survey Instrument
Title: *CityU Student Survey on Facial Recognition Technology and Biometric Privacy*
Sample Questions:

1. How many apps/platforms do you think collect your facial data?
    a. 0–2
    b. 3–5
    c. 6+
2. Do you believe companies delete your facial data after you stop using their services?

   a. Yes
   b. No
   c. Unsure
3. Rate your concern about unauthorized facial recognition surveillance (1 = Not Concerned, 5 = Extremely Concerned).

Demographics:

- Total Participants: 127
- Faculty Distribution: Engineering (32%), Business (28%), Humanities (22%), Science (18%)
- Recruitment Method: Campus posters, departmental emails.

Appendix B: Mock Datasets and Simulations
Dataset 1: Synthetic Facial Templates

- **Purpose**: Simulate unauthorized data retention in campus systems.
- **Structure**: 500 facial templates (PNG/JPG) with metadata (e.g., timestamp, location).
- **Source**: Generated using *Python's Faker Library* and *OpenCV face synthesis*.

Dataset 2: Simulated Breach Scenarios

- **Scenario 1**: *CampusSafe* third-party data sharing (200 templates).
- **Scenario 2**: Exam proctoring system adversarial attacks (150 AI-generated faces).